

High Performance and Distributed Computing for Big Data

Unit 3: Cloud Systems and Big Data Management

Session 5 - AWS Lambda

Francesc Solsona Tehas francesc.solsona@udl.cat

Universitat Rovira i Virgili and Universitat de Lleida

Cloud functions

Cloud functions enable **serverless** computing, allowing developers to run code without provisioning or managing servers.

Example

Automatically processing patient data uploads, triggering real-time alerts, and updating medical dashboards without infrastructure management.

How Serverless Functions Work



Figure 1: Serverless model



Figure 2: Serverless providers

AWS Lambda

Introduction to AWS Lambda

What is AWS Lambda?

- Event-driven, serverless computing service.
- Runs code in response to triggers.

Key benefits

- Automatic scaling and high availability.
- Pay-per-use billing model.

Common use cases

- Real-time file processing (e.g., medical imaging).
- Backend for web and mobile health applications.

Supported languages

- Python, Node.js, Java, Go, Ruby, .NET, and more.

AWS Lambda architecture overview



Photograph is taken



Amazon S3
Photo is uploaded
to an S3 Bucket



Lambda is
triggered



AWS Lambda
Lambda runs image
resizing code



Photo is resized into web,
mobile, and tablet sizes

Figure 3: Lambda Architecture

AWS Lambda architecture overview




Figure 4: Lambda Architecture

How AWS Lambda compares to EC2

AWS EC2

- Full control over infrastructure (turn it on/off, upgrading).
- Manual scalability management (want more? You have to add more).
- Fixed cost for uptime.

AWS Lambda

- No infrastructure management (serverless).
- Automatic scalability (from zero to thousands).
- Pay for actual execution time.

Use cases for AWS Lambda

- **Real-time data processing:** Lambda can analyze sensor data from wearable devices to identify irregular patterns.
- **Data aggregation:** Lambda can collect data from multiple clinical trial sites, aggregating information on patient outcomes, adverse events, and treatment efficacy preparing a dashboard for the researchers.
- **Data validation:** Lambda can validate lab results (such as blood tests) by checking for outliers or inconsistencies. For example, abnormal values outside the expected range may require further investigation. Alerts can be sent to the lab technician or the patient's healthcare provider.
- **Image processing:** Lambda can process images to identify patterns or anomalies.

Example: Counting cells at scale

Scenario

Imagine a lab that generates a large collection of cell images every time they run an experiment. Once the experiment is done, they want to know the number of cells in each image to analyze the results but they want this process to be automated and immediate.

Workflow

1. A lab uploads a collection of cell images to S3.
2. S3 events trigger Lambda functions (one event per image, one function per event).
3. Lambda functions process the images and count the cells.
4. Results are stored in S3.

Example: Counting cells at scale




Figure 5: Health Data Flow

Lab: Counting cells at scale

Pre-requisites

- A machine with AWS Credentials configured and Python 3.13 installed. (I am going to use the EC2 instance we created in the previous unit together with `uv` for managing python versions).
- The cell images downloaded and extracted. The image dataset for cell counting can be downloaded by clicking [here](#).

Goal

- Upload a collection of cell images to an S3 bucket and trigger a Lambda function to count the cells in each image. The lambda will store the results in another S3 bucket.

Steps

1. Create the buckets.
2. Create the Lambda function.
3. Add a trigger to the Lambda function.
4. Write the Lambda function code.
5. Create and publish a Lambda layer with the dependencies.
6. Upload the images to the input bucket.
7. Check the results in the output bucket and verify the Lambda logs.

Step 1: Create the buckets

As we did in the previous session, we are going to visit the S3 service in the AWS console and create two buckets, one for the input images and another for the output results. Leave everything as default and just set the name for each one as shown below:

- Input bucket: `medical-images-raw-[YOUR-NAME]`
- Output bucket: `medical-images-processed-[YOUR-NAME]`

In my case that will be `medical-images-raw-ferran-aran` and `medical-images-processed-ferran-aran`.

S3 Bucket names

Remember S3 bucket names must be unique across all AWS accounts. If you get an error when creating the bucket, try a different name (e.g., add a random number at the end).

Step 2: Create the Lambda function

We are going to search for `lambda` in the AWS Console as usual and click on the first result.





Figure 6: Lambda search

Step 2: Create the Lambda function

Now click on create function.



The screenshot shows the AWS Lambda console interface. At the top, there's a navigation bar with the AWS logo, a search bar, and various icons. Below the navigation bar, the left sidebar shows the 'Lambda' menu with options like 'Dashboard', 'Applications', and 'Functions'. The main content area is titled 'Functions (5)' and includes a search bar and a 'Create function' button highlighted with a red arrow. Below this, there's a table listing existing functions.

☐	Function name	Description	Package type	Runtime	Last modified
☐	RoleCreationFunction	Create SLR if absent	Zip	Python 3.9	13 minutes ago
☐	ModLabRole	updates LabRole to allow it to assume itself	Zip	Python 3.9	13 minutes ago
☐	MainMonitoringFunction	-	Zip	Python 3.9	13 minutes ago
☐	RedshiftEventSubscription	Create Redshift event subscription to SNS Topic.	Zip	Python 3.9	13 minutes ago

Figure 7: Create function

Step 2: Create the Lambda function

And fill the form like shown below (the function name doesn't matter but I suggest you use `count-cells`):

The screenshot shows the AWS Lambda 'Create function' console. The 'Author from scratch' tab is selected. The 'Basic information' section contains the following fields and options:


- Function name:** A text input field containing 'count-cells' (annotated with a red circle 1).
- Runtime:** A dropdown menu showing 'Python 3.13' (annotated with a red circle 3).
- Architecture:** Radio buttons for 'x86_64' (selected, annotated with a red circle 2) and 'arm64'.
- Permissions:** A section titled 'Change default execution role' with three radio buttons: 'Create a new role with basic Lambda permissions', 'Use an existing role' (selected, annotated with a red circle 4), and 'Create a new role from AWS policy templates'.
- Existing role:** A dropdown menu showing 'LabRole' (annotated with a red circle 6).

The 'Additional Configurations' section is partially visible at the bottom. At the bottom right, there are 'Cancel' and 'Create function' buttons, with the 'Create function' button annotated with a red circle 7.

Figure 8: Create function

Step 3: Add a trigger to the Lambda function

If we want our Lambda function to be executed when a new image is uploaded to the input bucket, we need to add a trigger. Click on the `+ Add trigger` button.



The screenshot displays the AWS Lambda console interface for a function named 'count-cells'. The top navigation bar includes the AWS logo, a search bar, and account information (United States (N. Virginia), Account ID: 7977-5604-8210, voclabs/user5018903=). The breadcrumb trail shows 'Lambda > Functions > count-cells'. The function name 'count-cells' is highlighted in a blue box. On the right, there are buttons for 'Throttle', 'Copy ARN', and 'Actions'. Below the function name, the 'Function overview' section is active, showing a 'Diagram' tab and a 'Template' tab. A diagram of the function is displayed, showing the 'count-cells' function icon and a 'Layers' section with '(0)' layers. A '+ Add trigger' button is visible. To the right, the 'Function ARN' is listed as 'arn:aws:lambda:us-east-1:797756048210:function:count-cells', the 'Description' is empty, and the 'Last modified' time is '2 minutes ago'. Below the overview, there are tabs for 'Code', 'Test', 'Monitor', 'Configuration', 'Aliases', and 'Versions'. The 'Code source' section is active, showing a search bar with 'count-cells' and buttons for 'Open in Visual Studio Code' and 'Upload from'. The bottom left shows the 'EXPLORER' sidebar with a file named 'lambda_function.py' selected.

Figure 9: Add trigger

Step 3: Add a trigger to the Lambda function

Start by searching for **S3** in the trigger configuration and then fill in the form like shown below:

The screenshot shows the 'Add trigger' configuration page in the AWS Lambda console. The page is titled 'Add trigger' and has a sub-header 'Trigger configuration' with an 'info' link. The configuration is divided into several sections: 'Bucket', 'Event types', 'Prefix - optional', 'Suffix - optional', and 'Recursive invocation'. The 'Bucket' section has a text input field containing 's3/medical-images-raw-ferran-aran' and a 'Bucket region' dropdown set to 'us-east-1'. The 'Event types' section has a dropdown menu set to 'All object create events'. The 'Prefix - optional' section has a text input field containing 'e.g. images/'. The 'Suffix - optional' section has a text input field containing 'e.g. .jpg'. The 'Recursive invocation' section has a checkbox labeled 'I acknowledge that using the same S3 bucket for both input and output is not recommended and that this configuration can cause recursive invocations, increased Lambda usage, and increased costs.' which is checked. At the bottom, there is a checkbox labeled 'Lambda will add the necessary permissions for AWS S3 to invoke your Lambda function from this trigger.' which is also checked. The page has a 'Cancel' button and an 'Add' button. Numbered red circles (1-5) and a red arrow point to specific elements: 1 points to the 'Trigger configuration' header, 2 points to the bucket name input, 3 points to the event types dropdown, 4 points to the 'All object create events' option, and 5 points to the 'I acknowledge...' checkbox.

Trigger configuration [info](#)

S3
aws asynchronous storage

Bucket
Choose or enter the ARN of an S3 bucket that serves as the event source. The bucket must be in the same region as the function.
s3/medical-images-raw-ferran-aran
Bucket region: us-east-1

Event types
Select the events that you want to have trigger the Lambda function. You can optionally set up a prefix or suffix for an event. However, for each bucket, individual events cannot have multiple configurations with overlapping prefixes or suffixes that could match the same object key.
All object create events

Prefix - optional
Enter a single optional prefix to limit the notifications to objects with keys that start with matching characters. Any special characters must be URL encoded.
e.g. images/

Suffix - optional
Enter a single optional suffix to limit the notifications to objects with keys that end with matching characters. Any special characters must be URL encoded.
e.g. .jpg

Recursive invocation
If your function writes objects to an S3 bucket, ensure that you are using different S3 buckets for input and output. Writing to the same bucket increases the risk of creating a recursive invocation, which can result in increased Lambda usage and increased costs. [Learn more](#)
☒ I acknowledge that using the same S3 bucket for both input and output is not recommended and that this configuration can cause recursive invocations, increased Lambda usage, and increased costs.

☒ Lambda will add the necessary permissions for AWS S3 to invoke your Lambda function from this trigger. [Learn more](#) about the Lambda permissions model.

[Cancel](#) [Add](#)

Figure 10: Add trigger

Step 3: Add a trigger to the lambda function

Okay so we've now configured our lambda to be triggered when a new object is created in the input bucket. But how do we access the image in the bucket from the lambda function?

Since we have configured the trigger to be an S3 event, AWS is going to send a **JSON object** to the lambda function with the information about the event.

Go back to the code tab as shown below.

The screenshot shows the AWS Lambda console for a function named 'count-cells'. At the top, a green status bar indicates that the trigger 'medical-images-raw-ferran-aram' has been successfully added. Below this, the 'Function overview' section is visible, featuring a diagram of the function's configuration. A red rectangular box highlights the 'S3' trigger in this diagram. To the right of the diagram, details such as 'Description', 'Last modified', 'Function ARN', and 'Function URL' are listed. At the bottom of the console, a navigation bar includes tabs for 'Code', 'Test', 'Monitor', 'Configuration', 'Aliases', and 'Versions'. A red circle containing the number '1' is positioned over the 'Code' tab, with a red arrow pointing to it. The 'Configuration' tab is currently selected, showing 'General configuration' details like 'Description', 'Memory', 'Ephemeral storage', 'Timeout', and 'Triggers'.

Figure 11: Code tab

Step 3: Add a trigger to the lambda function

Take a look at the default code that came with our lambda function:

```
import json

def lambda_handler(event, context):
    # TODO implement
    return {
        'statusCode': 200,
        'body': json.dumps('Hello from Lambda!')
    }
```

This is the basic structure of a lambda function. The `lambda_handler` function is the entry point of the lambda and it receives two arguments: `event` and `context`. The `event` argument is the JSON object we were talking about that contains the information about the event that triggered the lambda. So **anything we want to do with our lambda function has to be done inside this function.**

Step 3: Add a trigger to the lambda function

Let's see how the event looks like by printing it:

```
import json

def lambda_handler(event, context):
    print(json.dumps(event, indent=2))
    return {
        'statusCode': 200,
        'body': json.dumps('Hello from Lambda!')
    }
```

Step 3: Add a trigger to the lambda function

Once we are happy with the code, we need to “save” the changes by clicking on the **Deploy** button as shown below.




Figure 12: Deploy

Step 3: Add a trigger to the lambda function

Now we have to trigger the lambda function by uploading an image to the input bucket. You can do this by visiting the S3 service in the AWS Console and clicking on the input bucket

`medical-images-raw-[YOUR-NAME]` we created earlier. Then click on `Upload` and select an image to upload.

We can now go back our lambda, click on `Monitoring` and then on `View logs in CloudWatch` to see the logs of the lambda function.




Figure 13: CloudWatch logs

Step 3: Add a trigger to the lambda function

Click on the latest log stream to see the logs of the lambda function.

The screenshot shows the AWS CloudWatch console interface. The breadcrumb navigation at the top indicates the path: CloudWatch > Log groups > /aws/lambda/count-cells. The left-hand navigation pane lists various CloudWatch features, with 'Logs' and 'Log groups' being relevant to the current view. The main content area is titled '/aws/lambda/count-cells' and contains a 'Log group details' section. This section provides information about the log group, including its class (Standard), ARN, creation time (2 days ago), retention (Never expire), and stored bytes (4.31 KB). To the right of these details are sections for 'Metric filters', 'Subscription filters', 'Contributor Insights rules', 'KMS key ID', 'Anomaly detection', 'Data protection', 'Sensitive data count', 'Field indexes', and 'Transformer'. Below the details section is a horizontal tab bar with 'Log streams' selected. The 'Log streams (1)' section displays a single log stream with the name '2025/03/18/[\$LATEST]4b2d42ce5b5c41c09d0a55f0312c1e61'. A red arrow points to this log stream name. The log stream's last event time is shown as '2025-03-18 08:11:18 (UTC)'.

CloudWatch > Log groups > /aws/lambda/count-cells

/aws/lambda/count-cells

Log group details

- Log class** [Info](#)
Standard
- ARN**
[arn:aws:logs:us-east-1:281098262585:log-group:/aws/lambda/count-cells:*](#)
- Creation time**
2 days ago
- Retention**
Never expire
- Stored bytes**
4.31 KB
- Metric filters**
0
- Subscription filters**
0
- Contributor Insights rules**
-
- KMS key ID**
-
- Anomaly detection**
[Configure](#)
- Data protection**
-
- Sensitive data count**
-
- Field indexes**
[Configure](#)
- Transformer**
[Configure](#)

Log streams (1)

☐ Exact match ☐ Show expired [Info](#) < 1 > [Info](#)

☐	Log stream	Last event time
☐	<u>2025/03/18/[\$LATEST]4b2d42ce5b5c41c09d0a55f0312c1e61</u>	2025-03-18 08:11:18 (UTC)

Figure 14: CloudWatch logs

Step 3: Add a trigger to the lambda function

If everything went well you should see the event printed in the logs in the form of a JSON. There is lots of information but we are just interested in a couple of fields; the S3 bucket name and the object key.

2025-03-18T08:11:18.800Z	},
2025-03-18T08:11:18.800Z	"s3": {
2025-03-18T08:11:18.800Z	"s3SchemaVersion": "1.0",
2025-03-18T08:11:18.800Z	"configurationId": "f22958d6-15e6-4a93-b535-c2cd60c26f04",
2025-03-18T08:11:18.800Z	"bucket": {
2025-03-18T08:11:18.800Z	"name": "medical-images-raw-ferran-arann",
2025-03-18T08:11:18.800Z	"ownerIdentity": {
2025-03-18T08:11:18.800Z	"principalId": "AIE480VMJFCTZ"
2025-03-18T08:11:18.800Z	},
2025-03-18T08:11:18.800Z	"arn": "arn:aws:s3:::medical-images-raw-ferran-arann"
2025-03-18T08:11:18.800Z	},
2025-03-18T08:11:18.800Z	"object": {
2025-03-18T08:11:18.800Z	"key": "image-1.png",
2025-03-18T08:11:18.800Z	"size": 133675,
2025-03-18T08:11:18.800Z	"eTag": "e326aa977c8ba59ab4f3c943d0b1cb03",
2025-03-18T08:11:18.800Z	"sequencer": "0067D92AA399156004"
2025-03-18T08:11:18.800Z	}

Figure 15: CloudWatch logs

Step 3: Add a trigger to the lambda function

Okay now that we know we have the information we need to access the image in the S3 bucket, we can write some template code that accesses the image on the bucket that triggered the lambda and saves it to the `processed` bucket.

We'll be using the `boto3` library to interact with S3 as we did in the previous session. Go back to your lambda function on the "Code" tab and paste the following code. **Remember to change the bucket name** (note that the code takes 2 slides to fit):

```
import boto3
import json
import os
import urllib.parse

s3 = boto3.client('s3')
```

Step 3: Add a trigger to the lambda function

```
def lambda_handler(event, context):  
    # Extract bucket and image info from the S3 event  
    bucket = event['Records'][0]['s3']['bucket']['name']  
    key = urllib.parse.unquote_plus(event['Records'][0]['s3']['object']['key'])  
  
    # Download the image from raw S3  
    download_path = f'/tmp/{os.path.basename(key)}'  
    s3.download_file(bucket, key, download_path)  
  
    # Upload the image to processed S3 bucket  
    result_bucket = 'medical-images-processed-[YOUR-NAME]' # Replace with your bucket name  
    s3.upload_file(download_path, result_bucket, key + "-processed.png")  
  
    return {  
        'statusCode': 200,  
        'body': json.dumps(f"Processed {key}, found {cell_count} cells.")  
    }
```

Step 3: Add a trigger to the lambda function

Once again click on `Deploy` to save the changes, then go to the S3 bucket `medical-images-raw-[YOUR-NAME]` and upload an image to trigger the lambda function.

If everything went well you should see the same image uploaded to the `processed` bucket with the suffix `-processed.png` as shown below:




Figure 16: Processed image

Step 3: Add a trigger to the lambda function

Great so we now have a lambda function that:

1. Is triggered when an image is uploaded to the input bucket.
2. Downloads the image from the input bucket.
3. Uploads the image to the output bucket.

We are now going to design the code that processes the image to count the cells, and once we're happy with it we'll add it to the lambda function code.

Step 4: Write the Lambda function code

Lets open up the remote jupyter notebook on our EC2 instance that we have been using and create a new notebook. As a reminder, you can access the EC2 instance, activate the python environment (in this case we are using the sample `project2` environment we created in Session 3) and start the jupyter notebook server with the following commands:

```
ssh -i .ssh/aws-keypair ec2-user@<your-ec2-public-ip>  
cd project2  
source .project2-venv/bin/activate  
jupyter notebook --ip 0.0.0.0 --port 8888
```

Remember you can visit the second guide on the subject's website to see how to access the notebook server.

Step 4: Write the Lambda function code

With the cell images downloaded and extracted, we can now upload one of them to the notebook server from the browser to start working on the code for the Lambda function.

To upload an image to the notebook server, just drag and drop it to the browser window as shown below:




Figure 17: Upload image

Step 4: Write the Lambda function code

Next create a new notebook and paste the following code to install the dependencies.

```
!pip install matplotlib opencv-python  
!sudo dnf install mesa-libGL -y
```

And paste the following in another notebook to load the image and display it.

```
import cv2  
import matplotlib.pyplot as plt  
  
# Load the sample image  
image_path = 'image-1.png' # Replace with your image path  
image = cv2.imread(image_path)  
  
# Display the image  
plt.imshow(cv2.cvtColor(image, cv2.COLOR_BGR2RGB))  
plt.show()
```

Step 4: Write the Lambda function code

We are now free to work on whichever code we want to process the images. By using the Jupyter notebook we can test the code and see the results before deploying it to the Lambda function. For now, trust me and add the following code to the same notebook:

```
# Count cells by drawing contours around them
gray_image = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
adaptive_thresh = cv2.adaptiveThreshold(
    gray_image, 255, cv2.ADAPTIVE_THRESH_MEAN_C,
    cv2.THRESH_BINARY_INV, 65, 5
)
kernel = cv2.getStructuringElement(cv2.MORPH_ELLIPSE, (5,5))
morph_image = cv2.morphologyEx(adaptive_thresh, cv2.MORPH_OPEN, kernel, iterations=1)
contours, _ = cv2.findContours(
    morph_image, cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_SIMPLE
)

# Print the result
cell_count = len(contours)
print(f'Cell count: {cell_count}')
```

Step 4: Write the Lambda function code

Printing the result is fine but it would be even better if we could visualize the contours drawn around the cells. To do so, we can use the following code:

```
output_image = image.copy()
cv2.drawContours(output_image, contours, -1, (0, 255, 0), 2)

# Generate the images
fig, axes = plt.subplots(1, 2, figsize=(12, 6))

axes[0].imshow(cv2.cvtColor(image, cv2.COLOR_BGR2RGB))
axes[0].set_title('Original Image')
axes[0].axis('off')

axes[1].imshow(cv2.cvtColor(output_image, cv2.COLOR_BGR2RGB))
axes[1].set_title(f'Contours (Cells: {cell_count})')
axes[1].axis('off')

plt.show()
```

Step 4: Write the Lambda function code

The visualization should look like this:



Figure 18: Processed image

Step 4: Write the Lambda function code

By now our code does the following:

1. Load an image.
2. Process the image to count the cells.
3. Generate an image with the results.

We are now going to need to adapt this code to work in the Lambda function where it will have to read the image from the S3 bucket given its path and write the results back to another S3 bucket.

Step 4: Write the Lambda function code

In the AWS Console go to the Lambda service and click on the lambda function we created earlier, then scroll down to the code editor and paste the following code (note that the code takes 3 slides to fit):

```
import boto3
import cv2
import numpy as np
import json
import os
import urllib.parse

s3 = boto3.client('s3')

def lambda_handler(event, context):
    # Extract bucket and image info from the S3 event
    bucket = event['Records'][0]['s3']['bucket']['name']
    key = urllib.parse.unquote_plus(event['Records'][0]['s3']['object']['key'])
    original_name = os.path.splitext(os.path.basename(key))[0]

    download_path = f'/tmp/{os.path.basename(key)}'
```

Step 4: Write the Lambda function code

```
# Download and load the image from S3
s3.download_file(bucket, key, download_path)
image = cv2.imread(download_path)

# Count the cells using contours
gray_image = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
adaptive_thresh = cv2.adaptiveThreshold(
    gray_image, 255, cv2.ADAPTIVE_THRESH_MEAN_C,
    cv2.THRESH_BINARY_INV, 65, 5
)
kernel = cv2.getStructuringElement(cv2.MORPH_ELLIPSE, (5,5))
morph_image = cv2.morphologyEx(adaptive_thresh, cv2.MORPH_OPEN, kernel, iterations=1)
contours, _ = cv2.findContours(
    morph_image, cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_SIMPLE
)
cell_count = len(contours)
```

Step 4: Write the Lambda function code

```
# Draw contours on a copy of the image
output_image = image.copy()
cv2.drawContours(output_image, contours, -1, (0, 255, 0), 2)

# Save the processed image and upload it to the "processed" bucket
result_image_name = f"{original_name}-processed-{cell_count}-cells.png"
result_image_path = f'/tmp/{result_image_name}'
cv2.imwrite(result_image_path, output_image)
result_bucket = 'medical-images-processed-[YOUR-NAME]' # Replace with your bucket name
s3.upload_file(result_image_path, result_bucket, result_image_name)

return {
    'statusCode': 200,
    'body': json.dumps(f"Processed {key}, found {cell_count} cells.")
}
```

Step 4: Write the Lambda function code

Once the code is copied click on `Deploy` to save the changes.

The screenshot shows the AWS Lambda console interface for a function named 'count-cells'. The 'Code source' tab is active, displaying a Python file named 'lambda_function.py'. The code defines a 'lambda_handler' function that processes an image by finding contours and saving the result to an S3 bucket. In the left-hand 'EXPLORER' panel, the 'DEPLOY' section contains two buttons: 'Deploy (Ctrl+Shift+U)' and 'Test (Ctrl+Shift+I)'. A red arrow points to the 'Deploy' button. The breadcrumb navigation at the top shows 'Lambda > Functions > count-cells'.

```
10 def lambda_handler(event, context):
27     # ... (previous code) ...
28     kernel = cv2.getStructuringElement(cv2.MORPH_ELLIPSE, (5,5))
29     morph_image = cv2.morphologyEx(adaptive_thresh, cv2.MORPH_OPEN, kernel, iterations=1)
30     contours, _ = cv2.findContours(
31         morph_image, cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_SIMPLE
32     )
33     cell_count = len(contours)
34
35     # Draw contours on a copy of the image
36     output_image = image.copy()
37     cv2.drawContours(output_image, contours, -1, (0, 255, 0), 2)
38
39     # Save the processed image and upload it to the "processed" bucket
40     result_image_name = f'{original_name}-processed-{cell_count}-cells.png'
41     result_image_path = f'/tmp/{result_image_name}'
42     cv2.imwrite(result_image_path, output_image)
43     result_bucket = 'medical-images-processed-ferran-arann'
44     s3.upload_file(result_image_path, result_bucket, result_image_name)
45
```

Figure 19: Deploy

Step 5: Create and publish a Lambda layer with the dependencies

AWS Lambdas come with some python dependencies pre-installed such as `boto3` (which is the library we use to write and read from S3 buckets), but we are going to need to install `opencv-python` to process the images since it is not included by default.

To do so, we are going to create a Lambda layer with the dependencies and attach it to the Lambda function. Think of it as a way of packaging the needed dependencies so the lambda has them available when it runs.

We are going to need a machine with AWS CLI and its credentials configured as well as Python 3.13 and `zip` installed. I am going to use the EC2 instance we created in the previous unit together with `uv` for managing python versions. AWS CLI and `zip` are already installed in the instance.

Step 5: Create and publish a Lambda layer with the dependencies

Start by creating a folder which we'll use to build the layer and cd into it.

```
mkdir -p cell-count-layer/python/lib/python3.13/site-packages/  
cd cell-count-layer
```

Now create a virtual environment with  and install the dependencies.

```
uv venv --seed --python 3.13 .cell-count-venv  
source .cell-count-venv/bin/activate  
pip install opencv-python-headless -t python/lib/python3.13/site-packages
```

Step 5: Create and publish a Lambda layer with the dependencies

Now we are going to zip the contents of the folder to create the layer.

```
zip -r opencv.zip python
```

We'll need to create a bucket where we upload the layer so we can then import it to Lambda layers. You can use the AWS Console on your browser as we've done before or use the following command **where you have to replace [YOUR-NAME] with your name:**

```
aws s3 mb s3://layers-bucket-[YOUR-NAME]
```

Step 5: Create and publish a Lambda layer with the dependencies

Now publish the layer to the bucket.

```
aws s3 cp opencv.zip s3://layers-bucket-[YOUR-NAME]/
```


And finally, we are going to import the layer from the bucket to the Lambda layers.

```
aws lambda publish-layer-version \  
  --layer-name opencv \  
  --content S3Bucket=layers-bucket-[YOUR-NAME],S3Key=opencv.zip \  
  --compatible-runtimes python3.13
```

Let's now see how to add this layer to our lambda.

Step 5: Create and publish a Lambda layer with the dependencies

Visit the Lambda service on the AWS Console and look for the function we have been working on. Click on it.



The screenshot shows the AWS Lambda console interface. At the top, there's a navigation bar with the AWS logo, a search bar, and various utility icons. Below this, the breadcrumb navigation shows 'Lambda > Functions'. The main heading is 'Functions (6)', and it indicates 'Last fetched 2 minutes ago'. There's a search bar with the placeholder 'Filter by attributes or search by keyword'. To the right, there are 'Actions' and 'Create function' buttons. The functions are listed in a table with columns: Function name, Description, Package type, Runtime, and Last modified. The 'count-cells' function is highlighted with a red arrow.

☐	Function name	Description	Package type	Runtime	Last modified
☐	RoleCreationFunction	Create SLR if absent	Zip	Python 3.9	7 hours ago
☐	ModLabRole	updates LabRole to allow it to assume itself	Zip	Python 3.9	7 hours ago
☐	MainMonitoringFunction	-	Zip	Python 3.9	7 hours ago
☐	RedshiftEventSubscription	Create Redshift event subscription to SNS Topic.	Zip	Python 3.9	7 hours ago
☐	count-cells	-	Zip	Python 3.13	6 hours ago
☐	RedshiftOverwatch	Deletes Redshift Cluster if the count is more than 2.	Zip	Python 3.9	7 hours ago

Figure 20: Lambda layer

Step 5: Create and publish a Lambda layer with the dependencies

Click on the **Layers** section below the function's name.

The screenshot displays the AWS Lambda console for a function named 'count-cells'. The breadcrumb navigation shows 'Lambda > Functions > count-cells'. The function's name 'count-cells' is prominently displayed at the top left. Below it, the 'Layers' section is highlighted with a red arrow, indicating where to click to manage dependencies. The 'Layers' section currently shows '(0)' layers. To the left of the function diagram, there is an 'S3' trigger icon and a '+ Add trigger' button. To the right of the function diagram, there is a '+ Add destination' button. The right sidebar contains metadata: 'Description' (empty), 'Last modified' (1 second ago), 'Function ARN' (arn:aws:lambda:us-east-1:123456789012:function:count-cells), and 'Function URL' (Info). The bottom navigation bar includes tabs for 'Code', 'Test', 'Monitor', 'Configuration', 'Aliases', and 'Versions'. The 'Code' tab is currently selected, showing the 'Code source' section with an 'Info' link.

Step 5: Create and publish a Lambda layer with the dependencies

Click on **Add a layer**.

The screenshot shows the AWS Lambda console interface for a function named 'count-cells'. The breadcrumb navigation at the top indicates the path: Lambda > Functions > count-cells. The main content area is divided into three sections: Code properties, Runtime settings, and Layers.

Code properties

- Package size: 279 byte
- SHA256 hash: h2JxfYVX4TN1+LTCP3XyGf//SICnSXB4yekFDdhqYuE=
- Last modified: 22 seconds ago
- Encryption with AWS KMS customer managed KMS key

Runtime settings

- Runtime: Python 3.13
- Handler: lambda_function.lambda_handler
- Architecture: x86_64
- Buttons: Edit, Edit runtime management configuration

Layers

Merge order	Name	Layer version	Compatible runtimes	Compatible architectures	Version ARN
There is no data to display.					

Buttons: Edit, Add a layer (highlighted with a red arrow)

Figure 22: Lambda layer

Step 5: Create and publish a Lambda layer with the dependencies

Click on **Custom layers** and select the layer we just created. Finally click on **Add**.

The screenshot shows the 'Add layer' page in the AWS Lambda console. At the top, the breadcrumb 'Lambda > Add layer' is visible. The 'Function runtime settings' section shows 'Runtime' as 'Python 3.13' and 'Architecture' as 'x86_64'. The 'Choose a layer' section has three options: 'AWS layers', 'Custom layers' (selected), and 'Specify an ARN'. Below 'Custom layers', there is a list of layers. The first layer, 'opencv', is selected and highlighted with a red box and a red circle with the number '3'. Below it, the 'Version' is set to '1', also highlighted with a red box and a red circle with the number '5'. A red arrow with the number '1' points to the 'Custom layers' radio button. Another red arrow with the number '2' points to the 'Add' button at the bottom right. A third red arrow with the number '4' points to the 'Cancel' button at the bottom right. The 'Add' button is orange and labeled 'Add', while the 'Cancel' button is blue and labeled 'Cancel'.

Add layer

Function runtime settings

Runtime
Python 3.13

Architecture
x86_64

Choose a layer

Layer source [Info](#)
Choose from layers with a compatible runtime and instruction set architecture or specify the Amazon Resource Name (ARN) of a layer version. You can also [create a new layer](#).

☐ AWS layers
Choose a layer from a list of layers provided by AWS.

☒ Custom layers
Choose a layer from a list of layers created by your AWS account.

☐ Specify an ARN
Specify a layer by providing the ARN.

Custom layers
Layers created by your AWS account that are compatible with your function's runtime.

opencv

Version
1

Cancel Add

Figure 23: Lambda layer

Step 6: Upload the images to the input bucket

Remember the cell images can be found [here](#).

To upload them to the S3 bucket, we could do so by using the AWS Console as we did before, but this time we are going to use the AWS CLI to do it.

```
aws s3 cp ./cell_images s3://medical-images-raw-[YOUR-NAME]/ --recursive
```

The next slide contains a screenshot of the general steps to download, extract and upload the images to the S3 bucket.

Step 6: Upload the images to the input bucket

The screenshot shows a university course page for 'UNIVERSITAT ROVIRA I VIRGILI'. The page lists various documents and videos under 'Tema 3'. A red box highlights the document 'Cell images dataset (for Session 5 - AWS Lambda)' with a red arrow pointing to it labeled '1'. A Windows File Explorer window shows the 'Downloads' folder with a subfolder 'cell_images' highlighted by a red arrow labeled '2'. Inside 'cell_images', there are two files: 'cell_images.zip' and 'cell_images'. A red arrow labeled '3' points to 'cell_images'. A Windows PowerShell window shows the command 'ls .\cell_images\' and the output listing 10 image files (image-1.png to image-10.png) with a red arrow labeled '4' pointing to the command. Below the PowerShell window, a red box highlights the command 'ms s3 cp /cell_images s3://medical-images-raw-ferran-arann/ --recursive' with a red arrow labeled '5' pointing to it. The PowerShell window also shows the output of the command, listing the upload of each image file to the specified S3 bucket.

UNIVERSITAT ROVIRA I VIRGILI

Planificació Les meves aules Ajuda

Linux Parallelism 1.4 MB Document PDF

SLURM - Part I 3.9 MB Document PDF

SLURM Part I video 200.3 MB Fitxer de vídeo (MP4)

SLURM Part II 1.6 MB Document PDF

Tema 3

Session 1 & 2 - Cloud theory 1.8 MB Document PDF

Subject website with additional resources

Lab 01: Deploying your personal website

Session 3 - AWS EC2 (v2) 3.8 MB Document PDF

Session 4 - AWS S3 7.4 MB Document PDF

Cell images dataset (for Session 5 - AWS Lambda) 1.2 MB Fitxer (ZIP)

Microsoft T

Teams A_2C

Enregistrar

Guies docer

Guia 17705

Guia 17715

Calendari


Calendari com

Gestiona les us

Figure 24: Upload images

Step 6: Upload the images to the input bucket

If we check on the AWS Console the input bucket `medical-images-raw-[YOUR-NAME]` we should see the images uploaded



The screenshot shows the AWS S3 console interface. The left sidebar contains navigation links for Amazon S3, including General purpose buckets, Directory buckets, Table buckets, Access Grants, Access Points, Object Lambda Access Points, Multi-Region Access Points, Batch Operations, IAM Access Analyzer for S3, Storage Lens, Dashboards, Storage Lens groups, AWS Organizations settings, Feature spotlight, and AWS Marketplace for S3. The main content area displays the bucket 'medical-images-raw-ferran-arann' with tabs for Objects, Metadata, Properties, Permissions, Metrics, Management, and Access Points. The 'Objects' tab is active, showing a list of 10 objects. A red box highlights the first five objects: image-1.png, image-10.png, image-2.png, image-3.png, and image-4.png. The table lists the following objects:


Name	Type	Last modified	Size	Storage class
image-1.png	png	March 18, 2025, 08:42:58 (UTC+01:00)	130.5 KB	Standard
image-10.png	png	March 18, 2025, 08:42:58 (UTC+01:00)	124.9 KB	Standard
image-2.png	png	March 18, 2025, 08:42:58 (UTC+01:00)	125.0 KB	Standard
image-3.png	png	March 18, 2025, 08:42:58 (UTC+01:00)	123.9 KB	Standard
image-4.png	png	March 18, 2025, 08:42:58 (UTC+01:00)	128.9 KB	Standard
image-5.png	png	March 18, 2025, 08:42:58 (UTC+01:00)	124.0 KB	Standard
image-6.png	png	March 18, 2025, 08:42:58 (UTC+01:00)	121.5 KB	Standard
image-7.png	png	March 18, 2025, 08:42:58 (UTC+01:00)	123.3 KB	Standard
image-8.png	png	March 18, 2025, 08:42:58 (UTC+01:00)	123.9 KB	Standard
image-9.png	png	March 18, 2025, 08:42:58 (UTC+01:00)	123.5 KB	Standard

Figure 25: Uploaded images

Step 7: Check the results in the output bucket and verify the Lambda logs

If we've done everything correctly, a Lambda function should have been triggered for each image uploaded to the input bucket, and the processed images should be in the output bucket.

If we check S3 bucket `medical-images-processed-[YOUR-NAME]` we should see something like this:



The screenshot shows the Amazon S3 console interface for the bucket `medical-images-processed-ferran-arann`. The **Objects** tab is active, showing a list of 10 objects. A red box highlights the first nine objects, which are processed images. The tenth object is also visible. The table columns are Name, Type, Last modified, Size, and Storage class.

	Name	Type	Last modified	Size	Storage class
☐	image-1-processed-24-cells.png	png	March 18, 2025, 08:43:03 (UTC+01:00)	98.6 KB	Standard
☐	image-10-processed-20-cells.png	png	March 18, 2025, 08:43:04 (UTC+01:00)	93.5 KB	Standard
☐	image-2-processed-23-cells.png	png	March 18, 2025, 08:43:03 (UTC+01:00)	94.6 KB	Standard
☐	image-3-processed-19-cells.png	png	March 18, 2025, 08:43:03 (UTC+01:00)	93.9 KB	Standard
☐	image-4-processed-25-cells.png	png	March 18, 2025, 08:43:03 (UTC+01:00)	96.7 KB	Standard
☐	image-5-processed-27-cells.png	png	March 18, 2025, 08:43:03 (UTC+01:00)	94.3 KB	Standard
☐	image-6-processed-18-cells.png	png	March 18, 2025, 08:43:03 (UTC+01:00)	92.2 KB	Standard
☐	image-7-processed-20-cells.png	png	March 18, 2025, 08:43:03 (UTC+01:00)	93.5 KB	Standard
☐	image-8-processed-17-cells.png	png	March 18, 2025, 08:43:03 (UTC+01:00)	94.4 KB	Standard
☐	image-9-processed-22-cells.png	png	March 18, 2025, 08:43:03 (UTC+01:00)	95.0 KB	Standard

Figure 26: Processed images

Step 7: Check the results in the output bucket and verify the Lambda logs

And if we go back to our lambda and click on **Monitoring**. We should see a plot named **Invocations** that shows the number of times the lambda has been triggered.




Figure 27: Lambda invocations

Step 7: Check the results in the output bucket and verify the Lambda logs

We could also click on [View logs in CloudWatch](#) to see the logs of the lambda function as we did earlier.




Figure 28: Lambda logs

Recap

- AWS Lambda is a serverless computing service that runs code in response to events.
- Lambda functions are triggered by events from various sources, such as S3, API Gateway, and CloudWatch.
- Lambda functions can be used for real-time data processing, data aggregation, data validation, and image processing in healthcare applications.